

The Evolution of Computing


From the Perspective of the Cloud Service Provider

Fault Domains

Server Racks




Dedicated Servers



physical server


- A physical server is wholly utilized by a single customer.
- You have to guess your capacity, you'll overpay for an underutilized server
- Upgrading beyond your capacity will be slow and expensive
- Multiple apps can result in conflicts in resource sharing
- You have a guarantee of security, privacy and full utility of underlying resources

Virtual Machines




- You can run multiple Virtual Machines on one server
- A physical server can be shared by multiple customers
- You pay for a fraction of the server
- You'll overpay for an underutilized Virtual Machine
- Multiple apps on a single Virtual Machine can result in conflicts in resource sharing

Containers



- Virtual Machine running multiple containers
- You can maximum the utilize the available capacity which is more cost-effective
- Your containers share the same underlying OS so containers are more efficient than multiple VMs
- Multiple apps can run side by side without being limited to the same OS requirements and will not cause conflicts during resource sharing

Functions



- A managed VMs running managed containers running multiple functions.
- Also, known as Serverless Compute
- You upload a piece of code choose the amount of memory and duration.
- Only responsible for code and data, nothing else
- Very cost-effective, only pay for the time code is running, VMs only run when there is code to be executed
- Cold Starts is a side-effect of this setup

Azure Computing

Various Kinds

Types of Azure Computing



Azure Virtual Machines

Windows or Linux virtual machines (VMs). The most common type of compute.
You choose your OS, Memory, CPU, Storage. You share hardware with other customers



Azure Container Instances

Docker as a Service Run containerized apps on Azure without provisioning servers or VMs



Azure Kubernetes Service (AKS)

Kubernetes as a Service. Easy to deploy, manage and scale containerized applications.
Uses the open source Kubernetes (K8) software.



Azure Service Fabric

Tier-1 Enterprise Containers as a Service

Distributed systems platform. Runs in Azure or on-premises.
Easy to package, deploy, and manage scalable and reliable **microservices**.



Azure Functions

Event-driven, serverless compute (functions) run code without provisioning or managing servers.
You pay only for the compute time you consume.



Azure Batch

Plans, schedules and executes your batch computer workloads across running 100+ jobs in parallel.
Use Spot VMs to save money (previously used Low-priority VMs to save on compute)

Azure Virtual Machines



Azure VM Images



- A blueprint
- Defines
 - CPUs
 - GPUs
 - Storage
 - OS
 - Any pre-installed applications
 - Networking



Pricing Models

Pricing Model	Description	Best For
Pay-as-you-go	Pay only for compute time used (billed per second). No upfront cost or long-term commitment.	Short-term, unpredictable, or developmental workloads.
Reserved Instances (RIs)	Commit to a 1-year or 3-year term for up to 72% cost savings compared to pay-as-you-go.	Steady-state, predictable, continuous production workloads.
Spot VMs	Purchase unused Azure datacenter capacity at massive discounts (up to 90%). Caveat: Azure can evict and shut down your VM at any time if it needs the capacity back.	Interruptible workloads, batch processing, rendering, and large-scale testing.
Azure Hybrid Benefit	Bring your existing on-premises Windows Server or SQL Server licenses (with Software Assurance) to Azure to drastically reduce hourly VM costs.	Enterprises with existing Microsoft licensing agreements.




Azure VM Family Series

Category	VM Series & Specifics	Detailed Characteristics	Best Use Cases
General Purpose	B-Series (Burstable)	Cost-effective; earns "credits" during low usage to burst to full performance during spikes.	Small web servers, proof-of-concept (PoC) environments, idle workloads.
	D-Series	Enterprise standard; features faster processors and superior local disk performance.	Enterprise-grade apps, e-commerce systems, and standard databases.
	A-Series	Entry-level, lowest-cost virtual machines.	Dev/test environments and low-traffic servers.
Compute Optimized	F-Series	Highest CPU-to-memory ratio; prioritized calculation power over RAM volume.	Network appliances, batch processing, gaming servers, heavy logic web servers, and analytics.
Memory Optimized	E-Series	High RAM-to-CPU ratio for memory-intensive enterprise apps.	Memory-heavy enterprise applications.
	M-Series	The "heavyweights"; up to 11.4 TB of RAM and hundreds of vCPUs on a single machine.	Relational databases (SQL, Oracle), large in-memory caches (Redis), SAP HANA.
Storage Optimized	L-Series (and Ls)	High disk throughput/IOPS; features high-capacity, fast local NVMe storage attached to the host.	Big Data (NoSQL, Cassandra, MongoDB), large transactional databases with disk bottlenecks.
GPU / Accelerated	N-Series (NC, ND, NV)	Equipped with NVIDIA and AMD GPUs for specialized compute and visualization.	AI model training, deep learning, heavy video rendering, and remote 3D CAD workstations.
High Performance (HPC)	H-Series	Built with high-end CPUs and InfiniBand (high-bandwidth, low-latency network interfaces).	Scientific modeling, fluid dynamics, weather simulation, and financial risk modeling.




Azure VM Naming Convention



- **D** family is for General Purpose use cases, as we saw earlier.
- **C** sub-family stands for Confidential Computing. This means the VM uses hardware-based Trusted Execution Environments (like Intel SGX or AMD SEV-SNP) to encrypt data while it's being processed in memory.
- **a** feature stands for AMD-based processors. (Specifically, usually EPYC processors).
- **d** feature stands for Diskful. This indicates the VM comes with a local temporary disk (SSD) directly attached to the physical host.
- **s** feature means the VM is Premium Storage capable. It can be attached to Premium SSDs or Ultra Disks to achieve high IOPS and low latency.
- **v5** represents the 5th Generation. As Azure upgrades their data centers with newer, faster hardware (better CPUs, faster RAM, improved networking), they release new versions.

<https://azure.microsoft.com/en-in/pricing/details/virtual-machines/linux/#Linux>

Azure Virtual Machines Creation Process



- When you create a virtual machine, the following are also created
 - **Network Security Group (NSG):** A stateful firewall that controls inbound and outbound traffic via rules.
 - **Network Interface (NIC):** The bridge between the VM and the VNet - handles ip protocols and network communications.
 - **Virtual Network (VNet):** The isolated network boundary.
 - **Public IP Address** of the VNet
 - **OS Disk**
- Billed at an hourly rate
- 99.9% Availability (3-nines)
- You can attach multiple disks to the VM.
- OS and Security Patches are your responsibilities
- You deploy VMs across multiple AZs within a region, connected by low-latency fiber. This provides a 99.99% SLA.




OS Disk vs Data Disk


- OS Disk
 - The OS disk is automatically created and attached when you create a VM. It contains the boot volume and the pre-installed operating system (Windows or Linux) that you selected during deployment.
 - Primary Purpose: Booting the machine, housing the operating system registry/files, and running system-level applications.
 - It can be thought of as the C: drive in Windows or /dev/sda in Linux.
 - By default, Windows OS disks are usually 127 GB, and Linux OS disks are 30 GB. You can expand them, but there is a hard maximum of 4 TB.
 - Best Practice: You should never store application data, databases, or user files on the OS disk. If the OS crashes or requires a complete redeployment from an image, any custom data stored here could be lost or complicate the recovery process.
- Data Disk
 - Data disks are optional, user-created disks that you attach to your VM specifically to hold your application data.
 - Primary Purpose: Storing databases, application files, web server content, logs, and custom software.
 - Default Mount: You assign the drive letter or mount point (e.g., E:, F:, or /dev/sdc, /data).
 - Size and Scalability: A single data disk can be up to 32 TB. More importantly, you can attach multiple data disks to a single VM.
 - Performance: You can choose the exact performance tier for each individual disk based on your workload:
 - Standard HDD: Cheap, spinning disks for backup storage or rarely accessed data.
 - Standard SSD: Cost-effective solid-state drives for dev/test or light web servers.
 - Premium SSD: High-performance, low-latency drives for production databases.
 - Ultra Disk: Sub-millisecond latency for top-tier, extreme-IOPS enterprise databases.



Azure VM Lifecycle



Azure Scale Sets allow you to automatically increase or decrease your VM capacity



A Scale Set is a group **of identical VMs** (same Image and size)


Follow Along via Azure Portal

Azure Virtual Machines and Open AI
<https://github.com/SampurnR/temp2sql/>


Creating the Resource Group

Sign up and sign in are assumed

- In the search bar, search for resource groups, and click the first option.



- In the Resource Manager window, click on create.





Creating the Resource Group

- Enter the following values:
 - Subscription: Azure subscription 1 (or your subscription name)
 - Resource group name: sun-rg
 - Region: US East 2 (or the region of your own preference).
- Click on Review+create -> create

Basics

Tags

Review + create

Resource group - A container that holds related resources for an Azure solution. The resource group can include all the resources for the solution, or only those resources that you want to manage as a group. You decide how you want to allocate resources to resource groups based on what makes the most sense for your organization. [Learn more](#)

Subscription * ⓘ

Azure subscription 1

Resource group name * ⓘ



sun-rg

Region * ⓘ


(US) East US 2

Creating the Virtual Machine

- Search for and select Virtual machines.



- Click on create



Creating the Virtual Machine

- Select Virtual Machines

No virtual machine that runs Linux or Windows is currently selected.

**Virtual machine**
Best for lower-traffic workloads, testing, or to control or highly customize apps, OS, or file system. If your workload or traffic starts to grow, a VM can later be attached to a Virtual Machine Scale Set (VMSS).

**Virtual machine scale set (VMSS)**
Built-in scaling, performance optimization, load balancing, and batch management for 1 to 1,000 VMs (no added cost). Include multiple VM sizes, zones, regions, and domains, along with discounted Spot VMs.

**Presets**
Create a pre-configured VM designed to optimize memory, capacity, or for general use. Deploy as-is, or customize as needed.

**Hybrid, preconfigured, and high volume solutions**
Explore pre-configured Starter kits for Linux and Windows, Azure Arc hybrid infrastructure solutions, and more.



[Learn more about Windows virtual machines](#)

[Learn more about Linux virtual machines](#)



Creating the Virtual Machine

- Select the following options:
 - Subscription: Azure subscription 1 (or your subscription name)
 - Resource group name: sun-rg
 - Region: US East 2 (or the region of your own preference).
 - Availability options: No Infrastructure Redundancy required
 - Image: Ubuntu Server 24.04 or Ubuntu Server 22.04
 - VM Architecture: x64
 - Check “Run with Azure Spot Discount”. Let other spot-related options be default.

Subscription *	Azure subscription 1
Resource group *	sun-rg Create new
Instance details	
Virtual machine name *	sun
Region *	(US) East US 2 Deploy to an Azure Extended Zone
Availability options	No infrastructure redundancy required
Security type	Trusted launch virtual machines Configure security features
Image *	Ubuntu Server 24.04 LTS - x64 Gen2 See all images Configure VM generation
VM architecture	☐ Arm64 ☒ x64
Run with Azure Spot discount	☒
Eviction type	☒ Capacity only: Your virtual machine will be evicted when Azure's excess capacity disappears. ☐ Price or capacity: Your virtual machine will be evicted when Azure's excess capacity disappears, or costs exceed your specified max price.
Eviction policy	☒ Stop / Deallocate ☐ Delete



Creating the Virtual Machine

- Select the following options:
 - Size: Any with 8 GiB memory
 - Authentication type: Password
 - Set a username and password
 - Select inbound ports: HTTP(80) and SSH(22)

Size * ⓘ Standard_D2s_v3 - 2 vcpus, 8 GiB memory (\$0.03453) ▼
[See all sizes](#)
[View pricing history and compare prices in nearby regions](#)

Maximum price you want to pay per hour (USD) ⓘ 0.096
Enter a price greater than or equal to the hardware costs (\$0.03453)

Enable Hibernation ⓘ ☐
 ⓘ Hibernation does not currently support Trusted launch and Confidential virtual machines for Linux images. [Learn more](#) ⓘ

Administrator account

Authentication type ⓘ ☐ SSH public key
☒ Password

Username * ⓘ sun-text2sql ✓

Password * ✓

Confirm password * ✓

Inbound port rules

Select which virtual machine network ports are accessible from the public internet. You can specify more limited or granular network access on the Networking tab.

Public inbound ports * ⓘ ☐ None
☒ Allow selected ports

Select inbound ports * HTTP (80), SSH (22) ▼

⚠ This will allow all IP addresses to access your virtual machine. This is only recommended for testing. Use the Advanced controls in the Networking tab to



Creating the Virtual Machine

- Click “Review and Create”
- Click create
- Wait 1-5 minutes

Creating the Virtual Machine



CreateVm-canonical.ubuntu-24_04-lts-server-20260326131040 | Overview

Deployment

Search

Delete Cancel Redeploy Download Refresh

Overview

Inputs

Outputs

Template

✔ Your deployment is complete

Deployment name : CreateVm-canonical.ubuntu-24_04-lts-server-20260326131040

Subscription : Azure subscription 1

Resource group : sun-rg

Start time : 26/3/2026, 1:15:19 PM

Correlation ID : 1acf737b-03b3-4530-b29f-92f7b719c863

Deployment details

Resource	Type	Status	Operation details
✔ sun	Virtual machine	OK	Operation details
✔ sun678	Microsoft.Network/networkInter	OK	Operation details
✔ sun-vnet	Virtual network	OK	Operation details
✔ sun-nsg	Network security group	OK	Operation details
✔ sun-ip	Public IP address	OK	Operation details

Next steps

Set up auto-shutdown Recommended

Monitor VM health, performance, and network dependencies Recommended

Run a script inside the virtual machine Recommended

Go to resource

Create another VM

Scale out your VM

- Click Go to Resource



Setup the DNS

- In the Overview Page, under the networking header, click on “Not Configured” for DNS.

Operating system	: Linux (ubuntu 24.04)
Size	: Standard D2s v3 (2 vcpus, 8 GiB memory)
Primary NIC public IP	: 20.65.77.18 1 associated public IPs
Virtual network/subnet	: sun-vnet/default
DNS name	: Not configured
Health state	: -
Time created	: 26/3/2026, 7:45 AM UTC



Setup the DNS

- In the DNS name label, put in something you like, and click on Apply.

Summary

Assignment	Static
Public IP address ⓘ	20.65.77.18
Idle timeout (minutes) * ⓘ	4
DNS name label (optional)	text2sql .eastus2.cloudapp.azure.com

DNS alias records

Azure DNS alias records are qualifications on a DNS record set. Alias records can reference other Azure resources within your DNS zone. [Learn more.](#)

[+ Add](#) [↻ Refresh](#) [✕ Remove](#)

☐	Name	DNS zone
--------------------------	------	----------

No items found


Apply

Discard



Opening Gradio Port – 7860



- In the left tab, open Network settings within the Networking tab





Opening Gradio Port – 7860

- In the left tab, open Network settings within the Networking tab
- Click “Create port rule”



Opening Gradio Port – 7860

- Select Inbound port rule
- Enter the following values:
 - Port: 7860
 - Name: AllowGradio
 - Rest can be default

 **Add inbound security rule** ×
sun-nsg

Source ⓘ
Any

Source port ranges * ⓘ
*

Destination ⓘ
Any

Service ⓘ
Custom

Destination port ranges * ⓘ
7860 ✓

Protocol
☒ Any
☐ TCP
☐ UDP
☐ ICMPv4
☐ ICMPv6

Action
☒ Allow
☐ Deny

Priority * ⓘ
340

Name *
AllowGradi ✓

Description

Add Cancel

 Give feedback



Connecting to the VM

- From the Overview Networking tab, copy the public IP

Operating system	: Linux (ubuntu 24.04)
Size	: Standard D2s v3 (2 vcpus, 8 GiB memory)
Primary NIC public IP	: 20.65.77.18 1 associated public IPs
Virtual network/subnet	: sun-vnet/default
DNS name	: Not configured
Health state	: -
Time created	: 26/3/2026, 7:45 AM UTC



Connecting to the VM

- In your terminal, enter
`ssh username@public-ip`
- Enter “yes”
- Enter your password



Preparing VM for deployment

In your terminal, enter

```
sudo apt-get update
sudo apt-get install -y docker.io
sudo systemctl enable --now docker
sudo usermod -aG docker $USER
```

```
# Test your docker installation
exit
```

```
ssh username@public-ip
```

Enter your password

```
docker
```


```
sudo mkdir -p /app/text2sql
sudo chown -R sun-text2sql:sun-text2sql /app
```




**The following steps are optional. You can continue with your free Gemini API.
There is no free tier for Azure Open AI.**

Creating the Open AI Deployment

- Search for and select Azure OpenAI



- Click on create



Creating the Open AI Deployment

- Select Azure Open AI



No Azure OpenAI to display

Perform a wide variety of natural language tasks.

+ Create ▾



Foundry (Recommended)

Design, customize, and manage AI agents and models including OpenAI.



Azure OpenAI

Use Azure OpenAI models to create, test, deploy your AI solutions end-to-end.



Creating the Open AI Deployment

Enter the following values:

- Subscription: Azure subscription 1 (or your subscription name)
- Resource group name: sun-rg
- Region: US East 2 (or the region of your own preference).
- Name: sun-text2sql (this needs to be unique)
- Pricing Tier: Standard S0 (for beginners this is the only option - azure gives you discounts on the basis of your use. But the use has to be massive - billions of tokens.)
- Click on “Next” 3 times -> Review+create -> create.

Project Details

Subscription * ⓘ

Azure subscription 1

Resource group * ⓘ

sun-rg

[Create new](#)

i Azure AI services resource creation requires subscription registration, we detected that your selected subscription did not register Cognitive services resource type before, we will help you to register Cognitive services resource type when you select a subscription in subscription dropdown. Click to learn more how to check registration state for your selected subscription.

Instance Details

Region ⓘ

East US

Name * ⓘ

sun-text2sql

Pricing tier * ⓘ

Standard S0

[View full pricing details](#)



Creating the Open AI Deployment

Click on “Go to Resource”



Creating the Open AI Deployment

Click on “Go to Foundry portal”

A new tab will appear



Go to Foundry portal



Delete



Want to try the latest industry models and Agents ?

Azure OpenAI is limited to OpenAI models. Upgrading your resource type to Foundry

[Get Started](#)

^ Essentials

Resource group ([move](#)) : [sun-rg](#)

Status : Active

Location : East US

Subscription ([move](#)) : [Azure subscription 1](#)

Subscription ID : 5ad26ccf-1b33-41c2-90d9-b060ab1a7d61

Tags ([edit](#)) : [Add tags](#)

[Get Started](#)


Develop

Monitor




Creating the Open AI Deployment

Click on Deployments




Creating the Open AI Deployment

- In the Model deployments tab, click on Deploy model.



- Choose Deploy base model





<https://azure.microsoft.com/en-us/pricing/details/azure-openai/#pricing>

Go through the above link to decide which model you want to choose.

Recommendation: gpt-4.1



- Select a model

Models 69

Inference tasks

Choose a model to create a new deployment. For flows and other resources, create a deployment from their respective list. [Go to model catalog.](#)

Show description

Search

gpt-5-pro

Chat completion, Responses

gpt-5

Chat completion, Responses

gpt-4.1

Chat completion, Responses

gpt-4.1-mini

Chat completion, Responses

gpt-4o-transcribe-diarize

Speech to text

gpt-5-codex

Chat completion, Responses

Prev

Next

gpt-4.1

Task: Chat completion

Task: Responses

Direct from Azure models

Direct from Azure models are a select portfolio curated for their market-differentiated capabilities:

 - Secure and managed by Microsoft: Purchase and manage models directly through Azure with a single license, consistent support, and no third-party dependencies, backed by Azure's enterprise-grade infrastructure.
 - Streamlined operations: Benefit from unified billing, governance, and seamless PTU portability across models hosted on Azure - all part of Microsoft Foundry.
 - Future-ready flexibility: Access the latest models as they become available, and easily test, deploy, or switch between them within Microsoft Foundry; reducing integration effort.
 - Cost control and optimization: Scale on demand with pay-as-you-go flexibility or reserve PTUs for predictable performance and savings.

Learn more about [Direct from Azure models.](#)

Confirm

Cancel

Creating the Open AI Deployment

- Make sure you change deployment type to “Standard”
- Click on Deploy

Deploy gpt-4.1

Deployment name ^{*} 

Deployment type

Standard 

Standard: Pay per API call with lower rate limits. Adheres to Azure data residency promises. Best for intermittent workloads with low to medium volume. [Learn more about Standard deployments](#) .

Deployment details 

Model version upgrade policy

Upgrade once new default version becomes available 

Model version

2025-04-14 (Default) 

AI resource

sun-text2sql

 50K tokens per minute quota available for your deployment

Tokens per Minute Rate Limit 

50K

Corresponding requests per minute (RPM) = 50

Content filter 

DefaultV2 

Enable dynamic quota 

 Enabled

Deploy

Cancel



Creating the Open AI Deployment

- Make sure you copy the endpoint URI and key

Endpoint

Target URI

<https://sun-text2sql.openai.azure.com/openai/deployments/g...>



Key

.....





Github Actions

Your CI/CD Pipeline




What is Automation

Automation

Predictable set of events/actions

Movement of Data from one point to another





What is CICD

Continuous Integration (CI)

Automating the integration of code from multiple contributors. Tests and builds run on every commit to prevent breaking changes.

Continuous Deployment (CD)

The automated process of releasing validated code directly to production or a repository.

The Goal: Replace manual human intervention with consistent software execution.



Some notes about YAML

Configuration for most modern CI/CD tools, including GitHub Actions, is written in YAML.

YAML (a recursive acronym for "YAML Ain't Markup Language") is a human-readable data serialization language. **It is strictly used to represent data, not executable code.**



YAML Syntax

- **Indentation:** YAML uses whitespace (spaces) to denote structure and hierarchy. Tabs are strictly forbidden. Standard practice is to use two spaces per indentation level.
- **Key-Value Pairs:** Data is structured in dictionaries or hashes, represented by a key, a colon, a space, and the value.

```
name: GitHub Actions
version: 1.0
```
- **Lists (Arrays):** Denoted by a hyphen and a space (-). All items in a list must be at the same indentation level.

```
operating_systems:
  - ubuntu-latest
  - windows-latest
  - macos-latest
```
- **Booleans:** Can be represented as true/false.
- **Strings:** Generally do not require quotes unless they contain special characters (like :, {, }, [,], ,, &, *, #, ?, |, -, <, >, =, !, %, @, \).



The Anatomy of Github Actions

1. Workflows (Triggered by Events)

2. Jobs (Run in parallel by default)

3. Steps (Execute sequentially on one Runner)

4. Actions (Custom/reusable scripts)



The Anatomy of Github Actions

Workflows

Configurable automated processes (`github/workflows/*.yml`). Repos can have multiple.

Events

Activities that trigger workflows (e.g., `push`, `pull_request`, cron schedules).

Runners

The servers executing your workflows.

- GitHub-hosted: Ubuntu, Windows, macOS.
- Self-hosted: Your own configured servers.



The Anatomy of Github Actions

Jobs

- A set of steps executed on the same runner.
- Can be sequential (via dependencies) or parallel.

Steps

- Individual tasks (shell commands or actions).
- They share data and the workspace filesystem.

Actions

- Reusable applications for complex tasks (e.g., `actions/checkout@v4`).
- Generally sourced via GitHub Marketplace.



The Anatomy of Github Actions

Secrets

- Encrypted environment variables for sensitive data (API keys, passwords).
- `${{ secrets.SECRET_NAME }}`


Artifacts

- Files or directories generated during a run (compiled binaries, test reports).
- Usage: Saved for download or passed sequentially to subsequent jobs.



Adding Secrets to the Git Repo

- Click on Actions



- In the Security tab, click on Actions

Security

🔒 Advanced Security

🔑 Deploy keys

🔑 * Secrets and variables ^

Actions

Codespaces

Dependabot



Adding Secrets to the Git Repo

Add the following secrets:

- VM_HOST
- VM_USERNAME
- VM_PASSWORD
- AZURE_OPENAI_API_KEY
- AZURE_OPENAI_ENDPOINT
- AZURE_OPENAI_DEPLOYMENT_NAME
- AZURE_OPENAI_API_VERSION

Secrets

Variables

Environment secrets

This environment has no secrets.

Manage environment secrets

Repository secrets

This repository has no secrets.

New repository secret

After your tests

Remove anything with costs


Shut down the VM

- Click on Stop



Delete the Resource Group

- Click on Delete resource group



The screenshot shows the Azure portal interface for a resource group named 'sun-rg'. The 'Delete resource group' button is highlighted with a red box. The left sidebar shows the navigation menu with 'Overview' selected. The main content area displays the 'Essentials' section with details about the subscription and the 'Resources' section with a table of resources.

sun-rg Resource group

Search + Create ⚙ Manage view 🗑 Delete resource group 🔄 Refresh 📄 Export to CSV ⋮ Group by none 🔍 JSON View

Overview

- Activity log
- Access control (IAM)
- Tags
- Resource visualizer
- Events
- Settings
- Cost Management
- Monitoring
- Automation
- Help

Essentials

Subscription ([move](#)) Deployments
[Azure subscription 1](#) 1 Succeeded

Subscription ID Location
5ad26ccf-1b33-41c2-90d9-b060ab1a7d61 East US 2

Tags ([edit](#))
[Add tags](#)

Resources Recommendations

Filter for any field... Type equals all Location equals all + Add filter

☐	Name ↑	Type	Location
☐	sun	Virtual machine	East US 2
☐	sun-ip	Public IP address	East US 2
☐	sun-nsg	Network security group	East US 2
☐	sun-vnet	Virtual network	East US 2
☐	sun678	Network Interface	East US 2
☐	sun_OsDisk_1_90ff4f97937e4f26860a7349234a79e8	Disk	East US 2